

ASL_Project

May 11, 2022

1 Computer Vision for reading American Sign Language (ASL)

```
[ ]: from google.colab import drive
drive.mount('/content/drive')
```

1.1 Importing libraries

```
[ ]: import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np
import torch
from tqdm import tqdm
import torch.nn as nn
from torch import optim
import torch.nn.functional as F
import torchvision.datasets as datasets
from torchvision.io import read_image
from torchvision import transforms
from torchvision import models
from torch.utils.data import Dataset, DataLoader, \
    SubsetRandomSampler, random_split

import glob
from mpl_toolkits.axes_grid1 import ImageGrid
from skimage.transform import resize
```

```
[ ]: train_path='/content/drive/MyDrive/ASL_Project/dataset/asl_alphabet_train/'
    ↪ asl_alphabet_train'
test_path='/content/drive/MyDrive/ASL_Project/dataset/ASL_TEST/'
    ↪ asl_alphabet_test'
my_test_path='/content/drive/MyDrive/ASL_Project/My_Test_ASL'
```

1.1.1 Viewing the images from the training set

```
[ ]: train_folders=np.array(glob.glob('/content/drive/MyDrive/ASL_Project/dataset/
    ↳asl_alphabet_train/asl_alphabet_train/*'))

fig = plt.figure(figsize=(30, 30))
grid = ImageGrid(fig, 111,
                  nrows_ncols=(8, 8),
                  axes_pad=0,
                  )

l=0
for img in train_folders:
    it=np.array(glob.glob(img+'/*'))
    im=plt.imread(it[0])
    class_label=it[0].split('/')[8]
    grid[1].imshow(im,cmap='gray',interpolation='nearest')
    grid[1].text(10,25, class_label,fontsize=40,color='red')
    l+=1
```

1.1.2 Viewing the images from self-generated test data

```
[ ]: my_test_folders=np.array(glob.glob('/content/drive/MyDrive/ASL_Project/
    ↳My_Test_ASL/*'))

fig = plt.figure(figsize=(30, 30))
grid = ImageGrid(fig, 111,
                  nrows_ncols=(8, 8),
                  axes_pad=0,
                  )

l=0
for img in my_test_folders:
    it=np.array(glob.glob(img+'/*'))
    image=plt.imread(it[0])
    class_label=it[0].split('/')[6]
    image2=resize(image, (200, 200))
    grid[1].imshow(image2,cmap='gray',interpolation='nearest')
    grid[1].text(10,25, class_label,fontsize=40,color='red')
    l+=1
```

1.2 Transformations

Transformations on the Train dataset

```
[ ]: train_transforms = transforms.Compose([transforms.Resize((96,96)),
                                           transforms.RandomHorizontalFlip(p=0.3),
                                           transforms.ToTensor()])
```

Transformations on the Test dataset

```
[ ]: test_transforms = transforms.Compose([transforms.Resize((96,96)),
                                         transforms.ToTensor()])

[ ]: # performing resize and horizontal flip transformations on the train data set

data= datasets.ImageFolder(train_path, transform=train_transforms)

[ ]: # performing resize transformations on the two test data sets

test_data = datasets.ImageFolder(test_path, transform=test_transforms)
my_test_data=datasets.ImageFolder(my_test_path, transform=test_transforms)
```

1.2.1 Mapping of labels

```
[ ]: class_labels=data.class_to_idx
class_labels = {value:key for key, value in class_labels.items()}
print(class_labels)
print(f"Class to index mapping: {data.class_to_idx}")
```

1.3 Splitting of train and validation data set

```
[ ]: train_data_,val_data_ = random_split(data, (69600, 17400),
                                         generator=torch.Generator().manual_seed(25))
#splits the dataset randomly

[ ]: print('Total images in Train data set:', len(train_data_))
print('Total images in Validation data set:', len(val_data_))
print('Total images in Test data set:', len(test_data))
```

1.4 Loading images to data loaders

```
[ ]: train_loader = torch.utils.data.DataLoader(train_data_, batch_size=128,
        ↪shuffle=True)
val_loader = torch.utils.data.DataLoader(val_data_, batch_size=128,
        ↪shuffle=True)
test_loader = torch.utils.data.DataLoader(test_data)
my_test_loader = torch.utils.data.DataLoader(my_test_data)
```

1.5 Checking if GPU is available

```
[ ]: # use GPU to accelerate the training if available
device = "cuda" if torch.cuda.is_available() else "cpu"
print("Currently used device:", device)
```

1.6 Common functions for generating plots

```
[ ]: def plot_accuracy(train_accuracy, validation_accuracy):  
    '''  
    Function to plot the accuracy for Training and Validation sets  
  
    Parameters:  
    train_accuracy: list of accuracy for training set  
    validation_accuracy: list of accuracy for validation set  
  
    Return:  
    Plot of Train and Validation Accuracy  
  
    '''  
    # accuracy plots  
    plt.figure(figsize=(8,8))  
    plt.plot(train_accuracy, color='green', label='Train Accuracy')# plot train_  
↪accuracy  
    plt.plot(validation_accuracy, color='orange', label='Validataion_  
↪Accuracy')# plot validation accuracy  
    plt.xlabel('Epochs')  
    plt.ylabel('Accuracy')  
    plt.legend()  
    plt.show()
```

```
[ ]: def plot_loss(train_loss, validation_loss):  
    '''  
    Function to plot the accuracy for Training and Validation sets  
  
    Parameters:  
    train_loss: list of loss for training set  
    validation_loss: list of loss for validation set  
  
    Return:  
    Plot of Train and Validation Loss Curves  
  
    '''  
    plt.figure(figsize=(8,8))  
    plt.plot(train_loss, color='blue', label='Train Loss')# plot train loss  
    plt.plot(validation_loss, color='red', label='Validataion Loss')# plot_  
↪validation loss  
    plt.xlabel('Epochs')  
    plt.ylabel('Loss')  
    plt.legend()  
    plt.show()
```

1.7 Common function for performing training on the model

```
[ ]: def train_model(model, dataloader,optimizer,criterion):  
    '''  
    Function to train the model on training set of images  
  
    Parameters:  
    model: the trained model  
    dataloader: validation dataset  
    optimizer: optimizer with a learning rate  
    criterion: loss criterion  
  
    Return:  
    train_loss: train loss  
    train_accuracy: train accuracy  
  
    '''  
    print('Training the model')  
    model.train()  
    batch_loss, batch_accuracy = [], []  
  
    for k,(data,target) in tqdm(enumerate(dataloader),  
→total=int(len(train_loader)/train_loader.batch_size)):  
        data, target = data.to(device), target.to(device)  
  
        outputs = model(data)  
  
        loss = criterion(outputs, target)  
        acc = (outputs.argmax(1) == target).type(torch.float)  
        acc = torch.mean(acc)  
  
        optimizer.zero_grad()  
        loss.backward()  
        optimizer.step()  
  
        # Append to lists  
        batch_loss.append(loss.item())  
        batch_accuracy.append(acc.item())  
  
    train_loss = sum(batch_loss)/len(batch_loss)  
    train_accuracy = sum(batch_accuracy)/len(batch_accuracy)  
  
    print(f"Train Loss: {train_loss:.4f}, Train Accuracy: {train_accuracy:.2f}")  
  
    return train_loss, train_accuracy
```

1.8 Common function for performing validation on the model

```
[ ]: def validate_model(model, dataloader, criterion):  
    '''  
    Function to check performance on the validation dataset  
  
    Parameters:  
    model: the trained model  
    dataloader: validation dataset  
    criterion: loss criterion  
  
    Return:  
    val_loss: validation loss  
    val_accuracy: validation accuracy  
  
    '''  
  
    print('Validating the model')  
    model.eval()  
  
    with torch.no_grad():  
        for data, target in dataloader:  
            data, target = data.to(device), target.to(device)  
            outputs = model(data)  
            val_loss = criterion(outputs, target)  
            val_accuracy = (outputs.argmax(1) == target).type(torch.float).  
→mean()  
  
            print(f'Validation Loss: {val_loss.item():.4f}, Validation Accuracy:␣  
→{val_accuracy.item():.2f}')  
  
    return val_loss.item(), val_accuracy.item()
```

1.9 Common function for testing models

```
[ ]: def test_model(model, test_loader, x):  
    '''  
    Function to test model's performance on the test set. The function receives␣  
→two types of test_loaders.  
    One is the test set provided with the Kaggle dataset and the second one is␣  
→the dataset created by us.  
  
    Parameters:  
    model: the model to be used for testing  
    test_loader: contains the test set images  
    x: takes the value of 0 or 1. 0 indicates the test set sent to the function,␣  
→was the Kaggle
```

test data. 1 indicates the test set was the one created by us.

Return:

Prints accuracy on the test dataset

```
'''
true_labels=[] #correct labels
pred_labels=[] #labels which would be predicted by model
s=""

with torch.no_grad():
    total = 0
    correct = 0
    for image, labels in test_loader:
        model.eval()
        image=image.to(device)
        labels=labels.to(device)
        output=model(image)

        predictions=torch.max(output,1)[1]
        correct+=(predictions==labels).sum().cpu().numpy()
        total+=len(labels)

        labels=labels.cpu().detach().numpy()
        predictions=predictions.cpu().detach().numpy()

        true_labels.extend(labels)
        pred_labels.extend(predictions)

    acc = correct*100/total
    if(x==0):
        s='Kaggle'
    else:
        s='self generated'
    print('Accuracy on '+s+' test dataset: %f' %(acc))
```

1.9.1 Common Optimizer and Criterion functions

```
[ ]: def model_optimizer_criterion(model, learning_rate):
    '''
    Function to define the optimizer and criterion function for the model

    Parameters:
    learning_rate: learning rate for optimizer

    Return:
    optimizer: optimizer with a learning rate
```

```

        criterion: loss criterion
        '''
        # optimizer
        optimizer = torch.optim.Adam(model.parameters(), lr=learning_rate,
        ↪weight_decay = 0.0001)
        # loss function
        criterion = nn.CrossEntropyLoss()

        return optimizer, criterion

```

1.10 Deep Neural Network

1.10.1 Defining the NN class

```

[ ]: class NeuralNetwork(nn.Module):
    '''
        Defining a 4-layer Deep NN with Kaiming normal initialization and ReLU
        ↪activation functions
    '''

    def __init__(self):
        super(NeuralNetwork, self).__init__()

        self.flatten = nn.Flatten()
        #linear layers
        self.l1 = nn.Linear(3*96*96, 10000) # input image size is 96x96 and due
        ↪to RGB image there are 3 channels
        self.l2 = nn.Linear(10000, 2000)
        self.l3 = nn.Linear(2000, 250)
        self.l4 = nn.Linear(250, 29)

        nn.init.kaiming_normal_(self.l1.weight, mode='fan_in',
                                nonlinearity='relu')
        nn.init.kaiming_normal_(self.l2.weight, mode='fan_in',
                                nonlinearity='relu')
        nn.init.kaiming_normal_(self.l3.weight, mode='fan_in',
                                nonlinearity='relu')
        nn.init.kaiming_normal_(self.l4.weight, mode='fan_in',
                                nonlinearity='relu')

    def forward(self, x):
        #combining the entire model

        x = self.flatten(x)
        x = F.relu(self.l1(x))
        x = F.relu(self.l2(x))
        x = F.relu(self.l3(x))

```



```

        output = self.l4(x)

    return output

```

```

[ ]: nn_model = NeuralNetwork()
print('Layers of the updated pre-trained Neural Network model')
print(nn_model)

# trainable parameters
total_trainable_parameters = sum(para.numel() for para in nn_model.parameters()
    ↳if para.requires_grad)
print(f"{total_trainable_parameters:,} number of training parameters.")

```

```

[ ]: def neural_net_model():
    '''
        Function to train the model as well as check the performance on the
    ↳validation dataset. Performs early stopping when the
        validation loss exceeds the previous loss twice.
    '''

    num_epoch=40 #setting total number of epochs to be 40
    prev_loss=100 #initial loss value for early stopping
    max_penalty=3
    best_acc=0

    parameter_learning_rate=[0.001,0.0001,0.00001] # different values for
    ↳learning rate
    for i,learning_rate in enumerate(parameter_learning_rate): #Looping through
    ↳the learning rates
        train_loss , train_accuracy = [], []
        val_loss , val_accuracy = [], []
        counter=0
        model = NeuralNetwork().to(device)

        optimizer,criterion = model_optimizer_criterion(model, learning_rate)
        print("Training model on Learning Rate=",learning_rate)

        #model will train for a particular learning rate
        for epoch in range(num_epoch):# Looping through the epochs
            print(f"Epoch number {epoch+1} of {num_epoch}")

            #training the model
            train_epoch_loss, train_epoch_accuracy = train_model(model,
                                                                    train_loader,
                                                                    optimizer,
                                                                    criterion)

            #checking the performance on validation
            val_epoch_loss, val_epoch_accuracy = validate_model(model,

```

```

val_loader,
criterion)

    #append the loss and accuracy for every epoch for both train and
    ↪validation data set
    train_loss.append(train_epoch_loss)
    train_accuracy.append(train_epoch_accuracy)
    val_loss.append(val_epoch_loss)
    val_accuracy.append(val_epoch_accuracy)

    # Early Stopping
    if val_epoch_loss>prev_loss: # checking if validation loss for the
    ↪current epoch is greater than the previous loss
        counter+=1
        if counter>= max_penalty: # if the validation loss exceeds the
    ↪previous loss twice begin the early stopping process
            print('Early Stopping Initiated')
            break

    else:
        counter = 0

    prev_loss = val_epoch_loss # update the threshold loss to the
    ↪previous validation loss of the epoch

    #saving the best performing model
    if(val_epoch_accuracy>best_acc):
        best_acc=val_epoch_accuracy
        best_model=model

    #saving the model for a particular learning rate
    torch.save(model.state_dict(), '/content/drive/MyDrive/ASL_Project/
    ↪savedmodels/'+ 'model_neural_network'+str(learning_rate)+'.pth' )

    #plotting the accuracy and loss curves
    print('Plotting Validation and Training Accuracy Curves for Learning
    ↪Rate of ', learning_rate)
    plot_accuracy(train_accuracy, val_accuracy)
    print('Plotting Validation and Training Loss Curves for Learning Rate
    ↪of ', learning_rate)
    plot_loss(train_loss, val_loss)

    return best_model

```

1.10.2 Training the Neural Network model

```
[ ]: trained_neural_net_model=neural_net_model()
```

1.10.3 Testing the Neural Network model

Testing the model on given test dataset

```
[ ]: test_model(trained_neural_net_model,test_loader,0)
```

Testing the model on our test dataset

```
[ ]: test_model(trained_neural_net_model,my_test_loader,1)
```

1.11 Convolutional Neural Network(CNN)

```
[ ]: class CustomCNN(nn.Module):  
    ''' Creating a CNN model from scratch'''  
  
    def __init__(self):  
        super(CustomCNN, self).__init__()  
        self.conv1 = nn.Conv2d(3, 8, 5)  
        self.conv2 = nn.Conv2d(8, 16, 5)  
        self.conv3 = nn.Conv2d(16, 32, 3)  
        self.conv4 = nn.Conv2d(32, 64, 5)  
        self.fc1 = nn.Linear(64, 128)  
        self.fc2 = nn.Linear(128, 29)#output 29 classes  
        self.pool = nn.MaxPool2d(2, 2)  
  
    def forward(self, x):  
        x = self.pool(F.relu(self.conv1(x)))  
        x = self.pool(F.relu(self.conv2(x)))  
        x = self.pool(F.relu(self.conv3(x)))  
        x = self.pool(F.relu(self.conv4(x)))  
        bs, _, _, _ = x.shape  
        x = F.adaptive_avg_pool2d(x, 1).reshape(bs, -1)  
        x = F.relu(self.fc1(x))  
        x = self.fc2(x)  
        return x
```

```
[ ]: model_cnn = CustomCNN().to(device)  
print(model_cnn)  
# total parameters and trainable parameters  
total_params = sum(p.numel() for p in model_cnn.parameters())  
print(f"{total_params:,} total parameters.")  
total_trainable_params = sum(  
    p.numel() for p in model_cnn.parameters() if p.requires_grad)  
print(f"{total_trainable_params:,} training parameters.")
```

```
[ ]: def cnn_optimizer_criterion(model, learning_rate, wt_decay):
    '''
    Function to define the optimizer and criterion function for the CNN model

    Parameters:
    learning_rate: learning rate for optimizer

    Return:
    optimizer: optimizer with a learning rate
    criterion: loss criterion
    '''
    #optimizer
    optimizer = torch.optim.Adam(model.parameters()),
    lr=learning_rate, weight_decay=wt_decay)
    # loss function
    criterion = nn.CrossEntropyLoss()
    return optimizer, criterion
```

```
[ ]: def my_cnn_model():
    '''
    Function to train the model as well as check the performance on the
    validation dataset using CNN Architecture. Performs early stopping when the
    validation loss exceeds the previous loss thrice.
    '''
    num_epoch=40 #setting total number of epochs to be 40
    prev_loss=100 #initial loss value for early stopping
    max_penalty=3
    best_acc=0
    param_learning_rate=[0.1,0.01,0.001]
    param_weight_decay=[0.001,0.0001]
    # Hyper parameter tuning begins
    for i, learning_rate in enumerate(param_learning_rate): #Looping through
    different learning rates values
        for j, wt_decay in enumerate(param_weight_decay): #Looping through
    different weight decay values
            model = CustomCNN().to(device)

            optimizer, criterion=cnn_optimizer_criterion(model, learning_rate, wt_decay)
            train_loss , train_accuracy = [], []
            val_loss , val_accuracy = [], []
            for epoch in range(num_epoch):
                print(f"Epoch {epoch+1} of {num_epoch}")
                #training the model
                train_epoch_loss, train_epoch_accuracy = train_model(model,
                                                                    train_loader,
                                                                    optimizer,
                                                                    criterion)
```

```

        #checking the performance on validation
        val_epoch_loss, val_epoch_accuracy = validate_model(model,
                                                            val_loader,
                                                            criterion)

        # Early Stopping
        if val_epoch_loss>prev_loss: # checking if validation loss for
→the current epoch is greater than the previous loss
            counter+=1
            if counter>= max_penalty: # if the validation loss exceeds
→the previous loss twice begin the early stopping process
                print('Early Stopping Initiated')
                break

        else:
            counter = 0

        prev_loss = val_epoch_loss # update the threshold loss to the
→previous validation loss of the epoch

        #saving the best performing model
        if(val_epoch_accuracy>best_acc):
            best_acc=val_epoch_accuracy
            best_model=model

        #saving the model for a particular learning rate
        torch.save(model.state_dict(),'/content/drive/MyDrive/ASL_Project/
→savedmodels/'+ 'model_cnn_lr'+str(learning_rate)+'_wt_decay_'+str(wt_decay)+'.'
→pth' )

        #plotting the accuracy and loss curves
        print('Plotting Validation and Training Accuracy Curves for
→Learning Rate of ', learning_rate, ' and Weight Decay of ',wt_decay)
        plot_accuracy(train_accuracy,val_accuracy)
        print('Plotting Validation and Training Loss Curves for Learning
→Rate of ', learning_rate, ' and Weight Decay of ',wt_decay)
        plot_loss(train_loss,val_loss)

    return best_model

```

1.11.1 Training the CNN model

```
[ ]: trained_cnn_model=my_cnn_model()
```

1.11.2 Testing the CNN model

Testing the model on given test dataset

```
[ ]: test_model(trained_cnn_model, test_loader, 0)
```

Testing the model on our test dataset

```
[ ]: test_model(trained_cnn_model, my_test_loader, 1)
```

1.12 Resnet-50

1.12.1 Updating the fully connected layer of the pretrained Mobile Net V2 model

```
[ ]: def Resnet_50():  
    '''Using the concept of transfer learning on the pre-trained Resnet-50_  
→model and retraining the fully connected layer '''  
  
    model = models.resnet50(pretrained=True)  
  
    for param in model.parameters(): #freezing the pretrained_  
→layers  
        param.requires_grad = False  
  
    model.fc = nn.Sequential(nn.Linear(2048, 1000), #redefining the fully_  
→connected classifier layer  
                             nn.ReLU(),  
                             nn.Dropout(0.3),  
                             nn.Linear(1000, 500),  
                             nn.ReLU(),  
                             nn.Dropout(0.2),  
                             nn.Linear(500, 29))  
  
    return model  
  
[ ]: resnet_model = Resnet_50()  
print('Layers of the updated pre-trained Resnet 50 model')  
print(resnet_model)  
  
# total parameters  
total_parameters = sum(para.numel() for para in resnet_model.parameters())  
print(f"{total_parameters:,} number of total parameters.")  
  
# trainable parameters  
total_trainable_parameters = sum(para.numel() for para in resnet_model.  
→parameters() if para.requires_grad)  
print(f"{total_trainable_parameters:,} number of training parameters.")  
  
[ ]: def model_optimizer_criterion_res(model, learning_rate):  
    '''  
Function to define the optimizer and criterion function for the ResNet model
```

```

Parameters:
learning_rate: learning rate for optimizer

Return:
optimizer: optimizer with a learning rate
criterion: loss criterion
'''

# optimizer
optimizer = torch.optim.Adam(model.fc.parameters(), lr=learning_rate,
↪weight_decay = 0.0001)

# loss function
criterion = nn.CrossEntropyLoss()

return optimizer,criterion

```

```

[ ]: def resnet_50_model():
    '''
    Function to train the model as well as check the performance on the
    ↪validation dataset. Performs early stopping when the
    validation loss exceeds the previous loss twice.
    '''

    num_epoch=40 #setting total number of epochs to be 40
    prev_loss=100 #initial loss value for early stopping
    max_penalty=3
    best_acc=0

    parameter_learning_rate=[0.1,0.01,0.001] # different values for learning
    ↪rate
    for i,learning_rate in enumerate(parameter_learning_rate): #Looping through
    ↪the learning rates
        train_loss , train_accuracy = [], []
        val_loss , val_accuracy = [], []
        counter=0
        model = Resnet_50().to(device)

        optimizer,criterion = model_optimizer_criterion_res(model,
    ↪learning_rate)
        print("Training model on Learning Rate=",learning_rate)

        #model will train for a particular learning rate
        for epoch in range(num_epoch):# Looping through the epochs
            print(f"Epoch number {epoch+1} of {num_epoch}")

            #training the model
            train_epoch_loss, train_epoch_accuracy = train_model(model,
                                                                    train_loader,

```

```

optimizer,
criterion)

#checking the performance on validation
val_epoch_loss, val_epoch_accuracy = validate_model(model,
                                                    val_loader,
                                                    criterion)

#append the loss and accuracy for every epoch for both train and
→validation data set
train_loss.append(train_epoch_loss)
train_accuracy.append(train_epoch_accuracy)
val_loss.append(val_epoch_loss)
val_accuracy.append(val_epoch_accuracy)

# Early Stopping
if val_epoch_loss>prev_loss: # checking if validation loss for the
→current epoch is greater than the previous loss
    counter+=1
    if counter>= max_penalty: # if the validation loss exceeds the
→previous loss twice begin the early stopping process
        print('Early Stopping Initiated')
        break

    else:
        counter = 0

    prev_loss = val_epoch_loss # update the threshold loss to the
→previous validation loss of the epoch

#saving the best performing model
if(val_epoch_accuracy>best_acc):
    best_acc=val_epoch_accuracy
    best_model=model

#saving the model for a particular learning rate
torch.save(model.state_dict(), '/content/drive/MyDrive/ASL_Project/
→savedmodels/'+ 'model_mobile_net'+str(learning_rate)+'.pth' )

#plotting the accuracy and loss curves
print('Plotting Validation and Training Accuracy Curves for Learning
→Rate of ', learning_rate)
plot_accuracy(train_accuracy, val_accuracy)
print('Plotting Validation and Training Loss Curves for Learning Rate
→of ', learning_rate)
plot_loss(train_loss, val_loss)

return best_model

```


1.12.2 Training the Resnet 50 model

```
[ ]: trained_resnet_50_model = resnet_50_model()
```

1.12.3 Testing the Resnet 50 model

Testing the model on given test dataset

```
[ ]: test_model(trained_resnet_50_model, test_loader, 0)
```

Testing the model on our test dataset

```
[ ]: test_model(trained_neural_net_model, my_test_loader, 1)
```

1.13 Mobile-Net V2

1.13.1 Updating the pretrained Mobile Net V2 model

```
[ ]: def ModelMob():  
    '''Using the concept of transfer learning on the pre-trained Mobile Net V2_  
    ↪model and retraining the fully connected layer'''  
  
    model = models.mobilenet_v2(pretrained=True)  
  
    for param in model.parameters():  
        #freezing the_  
    ↪pretrained layers  
        param.requires_grad = False  
  
    model.classifier = nn.Sequential(nn.Linear(1280, 128),    #redefining the_  
    ↪fully connected classifier layer  
                                     nn.ReLU(),  
                                     nn.Dropout(0.2),  
                                     nn.Linear(128, 64),  
                                     nn.ReLU(),  
                                     nn.Dropout(0.2),  
                                     nn.Linear(64, 29))  
  
    return model
```

```
[ ]: mob_model = ModelMob()  
print('Layers of the updated pre-trained Mobile Net-V2 model')  
print(mob_model)  
# total parameters  
total_parameters = sum(para.numel() for para in mob_model.parameters())  
print(f"{total_parameters:,} number of total parameters.")  
  
# trainable parameters  
total_trainable_parameters = sum(para.numel() for para in mob_model.  
    ↪parameters() if para.requires_grad)
```

```
print(f"{total_trainable_parameters:}, number of training parameters.")
```

1.13.2 Training the Mobile Net V2 model

```
[ ]: def mobile_net_model():
    '''
    Function to train the model as well as check the performance on the
    ↪ validation dataset. Performs early stopping when the
    validation loss exceeds the previous loss twice.
    '''

    num_epoch=40 #setting total number of epochs to be 40
    prev_loss=100 # initial loss value for early stopping
    max_penalty=2
    best_acc=0

    parameter_learning_rate=[0.1,0.01,0.001]# different values for learning
    ↪ rate
    for i,learning_rate in enumerate(parameter_learning_rate):#Looping through
    ↪ the learning rates
        train_loss , train_accuracy = [], []
        val_loss , val_accuracy = [], []
        counter=0
        model = ModelMob().to(device)
        optimizer,criterion=model_optimizer_criterion(model, learning_rate)
        print("Training model on Learning Rate=",learning_rate)
        #model will train for a particular learning rate
        for epoch in range(num_epoch):# Looping through the epochs
            print(f"Epoch number {epoch+1} of {num_epoch}")
            #training the model
            train_epoch_loss, train_epoch_accuracy = train_model(model,
                                                                    train_loader,
                                                                    optimizer,
                                                                    criterion)

            #checking the performance on validation
            val_epoch_loss, val_epoch_accuracy = validate_model(model,
                                                                val_loader,
                                                                criterion)

            #append the loss and accuracy for every epoch for both train and
            ↪ validation data set
            train_loss.append(train_epoch_loss)
            train_accuracy.append(train_epoch_accuracy)
            val_loss.append(val_epoch_loss)
            val_accuracy.append(val_epoch_accuracy)
            # Early Stopping
            if val_epoch_loss>prev_loss:# checking if validation loss is
            ↪ greater than the previous loss
```

```

        counter+=1
        if counter>= max_penalty:# if the validation loss exceeds the
→threshold loss twice begin the early stopping process
            print('Early Stopping Initiated')
            break
    else:
        counter = 0

    prev_loss=val_epoch_loss# update the previous loss to the previous
→validation loss of the epoch

    #saving the best performing model
    if(val_epoch_accuracy>best_acc):
        best_acc=val_epoch_accuracy
        best_model=model
    #saving the model for a particular learning rate
    torch.save(model.state_dict(),'/content/drive/MyDrive/ASL_Project/
→savedmodels/'+ 'model_mobile_net'+str(learning_rate)+'.pth' )
    #plotting the accuracy and loss curves
    print('Plotting Validation and Training Accuracy Curves for Learning
→Rate of ', learning_rate)
    plot_accuracy(train_accuracy,val_accuracy)
    print('Plotting Validation and Training Loss Curves for Learning Rate
→of ', learning_rate)
    plot_loss(train_loss,val_loss)

    return best_model

```

```
[ ]: trained_mobile_net_model=mobile_net_model()
```

1.13.3 Testing the Mobile Net V2 model

Testing the model on given test dataset

```
[ ]: test_model(trained_mobile_net_model,test_loader,0)
```

Testing the model on our test dataset

```
[ ]: test_model(trained_mobile_net_model,my_test_loader,1)
```